

Software Requirements Specification for Software Engineering: Code Plagiarism Detector

Team 2, SyntaxSentinals
Mohammad Mohsin Khan
Lucas Chen
Dennis Fong
Julian Cecchini
Luigi Quattrociochi

October 11, 2024

Contents

1	Purpose of the Project	vi
1.1	User Business (Cases)	vi
1.2	Goals of the Project	vii
2	Stakeholders	ix
2.1	Client	ix
2.2	Customer	x
2.3	Other Stakeholders	x
2.4	Hands-On Users of the Project	x
2.5	Personas	x
2.6	Priorities Assigned to Users	xi
2.7	User Participation	xi
2.8	Maintenance Users and Service Technicians	xi
3	Mandated Constraints	xi
3.1	Solution Constraints	xi
3.2	Implementation Environment of the Current System	xi
3.3	Partner or Collaborative Applications	xi
3.4	Off-the-Shelf Software	xii
3.5	Anticipated Workplace Environment	xii
3.6	Schedule Constraints	xiii
3.7	Budget Constraints	xiii
3.8	Enterprise Constraints	xiii
4	Naming Conventions and Terminology	xiii
4.1	Glossary of All Terms, Including Acronyms, Used by Stakeholders involved in the Project	xiii
5	Relevant Facts And Assumptions	xiv
5.1	Relevant Facts	xiv
5.2	Business Rules	xiv
5.3	Assumptions	xv
6	The Scope of the Work	xv
6.1	The Current Situation	xv
6.2	The Context of the Work	xv
6.3	Work Partitioning	xvi

6.4	Specifying a Business Use Case (BUC)	xvi
7	Business Data Model and Data Dictionary	xvii
7.1	Business Data Model	xvii
7.2	Data Dictionary	xvii
8	The Scope of the Product	xviii
8.1	Product Boundary	xviii
8.2	Product Use Case Table	xviii
8.3	Individual Product Use Cases (PUC's)	xviii
9	Functional Requirements	xviii
9.1	Functional Requirements	xviii
10	Look and Feel Requirements	xix
10.1	Appearance Requirements	xix
10.2	Style Requirements	xix
11	Usability and Humanity Requirements	xx
11.1	Ease of Use Requirements	xx
11.2	Personalization and Internationalization Requirements	xx
11.3	Learning Requirements	xxi
11.4	Understandability and Politeness Requirements	xxi
11.5	Accessibility Requirements	xxi
12	Performance Requirements	xxi
12.1	Speed and Latency Requirements	xxi
12.2	Safety-Critical Requirements	xxii
12.3	Precision or Accuracy Requirements	xxii
12.4	Robustness or Fault-Tolerance Requirements	xxii
12.5	Capacity Requirements	xxii
12.6	Scalability or Extensibility Requirements	xxii
12.7	Longevity Requirements	xxiii
13	Operational and Environmental Requirements	xxiii
13.1	Expected Physical Environment	xxiii
13.2	Wider Environment Requirements	xxiii
13.3	Requirements for Interfacing with Adjacent Systems	xxiii
13.4	Productization Requirements	xxiii

13.5 Release Requirements	xxiii
14 Maintainability and Support Requirements	xxiii
14.1 Maintenance Requirements	xxiii
14.2 Supportability Requirements	xxiv
14.3 Adaptability Requirements	xxv
15 Security Requirements	xxvi
15.1 Access Requirements	xxvi
15.2 Integrity Requirements	xxvi
15.3 Privacy Requirements	xxvi
15.4 Audit Requirements	xxvi
15.5 Immunity Requirements	xxvii
16 Cultural Requirements	xxvii
16.1 Cultural Requirements	xxvii
17 Compliance Requirements	xxvii
17.1 Legal Requirements	xxviii
17.2 Standards Compliance Requirements	xxix
18 Open Issues	xxxi
18.1 System Design Uncertainty	xxxi
18.2 Data Collection Challenges	xxxi
18.3 Scalability Concerns	xxxi
19 Off-the-Shelf Solutions	xxxii
19.1 Ready-Made Products	xxxii
19.2 Reusable Components	xxxii
19.3 Products That Can Be Copied	xxxii
20 New Problems	xxxiii
20.1 Effects on the Current Environment	xxxiii
20.2 Effects on the Installed Systems	xxxiii
20.3 Potential User Problems	xxxiii
20.4 Limitations in the Anticipated Implementation Environment That May Inhibit the New Product	xxxiv
20.5 Follow-Up Problems	xxxiv

21 Tasks	xxxiv
21.1 Project Planning	xxxiv
21.2 Planning of the Development Phases	xxxvii
22 Migration to the New Product	xxxix
22.1 Requirements for Migration to the New Product	xxxix
22.2 Data That Has to be Modified or Translated for the New System	xl
23 Costs	xl
24 User Documentation and Training	xli
24.1 User Documentation Requirements	xli
24.2 Training Requirements	xli
25 Waiting Room	xli
26 Ideas for Solution	xli

Revision History

Date	Version	Notes
October 9	1.0	First iteration of complete document

1 Purpose of the Project

1.1 User Business (Cases)

The purpose of the project is first and foremost illustrated through the fulfillment of use cases our stakeholders are interested. These use cases and associated stakeholder interest are stated below.

1. Plagiarism Assessment

- 1.1 user selects option from main UI screen to provide dataset for scanning
- 1.2 user is prompted to provide dataset
 - 1.2.1 user inserts data from their computer
 - 1.2.2 user cancels and is sent back to main UI screen (1.1 continues)
- 1.3 user initiates analysis and progress indicator appears
 - 1.3.1 error occurs during analysis and user is prompted to restart or cancels
 - 1.3.1.1 user selects to restart analysis (1.3 continues)
 - 1.3.1.2 user cancels and returns to main UI (1.1 continues)
 - 1.3.2 progress bar completes and success message is displayed (proceeds to 1.4)
- 1.4 system provides results in a UI display accompanied by option to download result or return to main UI
 - 1.4.1 user selects download option
 - 1.4.1.1 download fails accompanied by error message (continues from 1.4)
 - 1.4.1.2 download succeeds and system provides dialogue to check for download after which user returns to main UI
- 1.5 system has provided result of analysis

Stakeholder Interest: meets desire to gain direction on whether or not plagiarism has occurred within their environment

2. Plagiarism Detector Adjustment

- 2.1 user selects option from main UI to open settings for app
- 2.2 user can adjust threshold used for detection, and possibly other components such as version of model architecture used (not confirmed at the moment)
- 2.3 user saves settings through UI provided option
 - 2.3.1 System displays setting changes are saved, displaying success message, prompting whether user desires to continue or not
 - 2.3.1.1 user selects to continue (2.2 continues)
 - 2.3.1.2 user selects to finish (2.4 continues)
 - 2.3.2 System displays setting changes did not save and provides dialogue to determine error, giving option to try again or return to main UI
 - 2.3.2.1 user selects to try again (2.2 continues)
 - 2.3.2.2 user selects to return to main UI (2.1 continues)
- 2.4 system maintains settings and applies them for functions until application restart

Stakeholder Interest: allows all stakeholders to cater app more towards their particular problem.

1.2 Goals of the Project

Goal	Explanation	Reason
Ease of Use	Detector has an intuitive way to insert data and obtain results	This application is expected to be used as a secondary tool for teachers/professors when administering assignments. It should not require in-depth learning, or it will be too inconvenient as an assistant tool for detecting plagiarism. (Measured by actions to complete analysis)

Clarity of Output	Detector explains how to interpret outputs clearly, leaving no ambiguity in whether plagiarism is suspected	If the user does not comprehend the output, it may result in unjust accusations or undetected plagiarism. (Measured by lines of description or number of users who correctly interpret output)
Real-Time Processing	The detector computes results on a dataset of code snippets quickly, enabling professors to incorporate them into evaluations	Since multiple assignments are administered over several weeks, the detector must be fast enough to be realistic for daily use. (Measured by execution time)
False Positive Accuracy	The detector prioritizes minimizing false positives over false negatives	In this case, a false positive could cause harm to an innocent student, while a false negative allows a violation to go unnoticed. The focus is on protecting innocent students. (Measured by false positives and negatives using recall, precision, etc.)
Ethically Sourced Data	The detector uses only data that openly discloses its origins and all data used for the detector is stated for all to see.	In modern ML development, it has become a hot topic for how models get their data. This is because many modern models have used datasets that contain information that was taken from individuals without consent (such as pieces of art). It is important to our team to make it clear this is not precedent and that individuals should have clear consent in being used for training models. (Measured by having accreditation for datasets involved in training)

Stretch Goal	Explanation	Reason
Online Learning	Provide ability for user to train model on their own datasets.	Without a sufficiently large amount of data to train on (more than will be seen over the duration of 8 months), the model will be biased to a degree and not as widely applicable to different sets of code. If the detector is given the ability to be trained by the user, they can better customize it for their own needs (measure is whether or not the user can change conduct training)
Language Agnostic	The detector can analyze code from a multitude of languages	Having a detector that can draw patterns across different languages will make it adoptable by a wider set of professors who may conduct courses in less popular languages that our detector may have not dealt with at all before. If the detector is restrained to languages such as python or java, we will alienate some of our primary stakeholders.

2 Stakeholders

2.1 Client

The primary clients for this project are the computing and software departments of academic institutions, and code competition administrators. These clients seek an advanced plagiarism detection system that overcomes the limitations of existing tools like MOSS. In academic settings, the system will help maintain academic integrity, while in code competitions, it will ensure fair play by preventing plagiarism among participants.

2.2 Customer

The primary customers for this project are professors, course instructors, and code competition organizers. Professors and instructors will use the system to evaluate student submissions to identify plagiarism, while competition organizers will ensure that all participants submit original code.

2.3 Other Stakeholders

Other stakeholders include:

- Students: Indirectly affected, as their work will be evaluated by this system, and it is important to ensure that students don't get falsely accused of plagiarism.
- Competition Participants: In coding competitions, the participants rely on the system to ensure the competition they participate in is fair.

2.4 Hands-On Users of the Project

The hands-on users are the professors, teaching assistants, and code competition organizers who will directly interact with the system. They will use it to upload code submissions, compare entries, and review plagiarism reports.

2.5 Personas

- Professor: Dr. Onjama Wembo - A computer science professor who frequently assigns coding tasks and reviews student submissions.
- Student: John Johnson - An honest computer science student who expects the system to verify their work as valid and un plagiarised.
- Competition Organizer: Sung Yuhee - A competition organizer who uses the system to ensure participants submit original work to ensure fair play in the competition.
- Competition Participant: Giorno Capio - A competition participant aiming for a top score in the competition, who is hoping the system does not misclassify their work with someone elses who may have similar code.

2.6 Priorities Assigned to Users

- High Priority: Professors, competition organizers, and instructors - They rely on the system to evaluate submissions to ensure integrity
- Low Priority: Students and competition participants - They are not the direct users of the system but rely on it for correct evaluations

2.7 User Participation

User participation is essential for the development and testing of the system. Professors will provide feedback during development and testing to ensure the system meets their needs. Regular feedback will help improve the system's accuracy and reliability.

2.8 Maintenance Users and Service Technicians

If possible, system administrators and IT staff will maintain the system. They will troubleshoot issues to keep the system functional, as well as oversee updates to the system.

3 Mandated Constraints

3.1 Solution Constraints

Constraint	Rationale
Our product shall follow a zero data retention policy.	We should comply with privacy laws regarding storage of sensitive information.

3.2 Implementation Environment of the Current System

3.3 Partner or Collaborative Applications

Our project will use Google Colab to train and tune our language models for the NLP component of our tool.

Constraint	Rationale
The application code should run on both Windows and Linux.	Users of our product who wish to host their own instance will be using a variety of server hardware providers, including Microsoft Azure and Amazon AWS.
The user interface should be accessible via a web browser on a computer with an internet connection.	Users of our product will access the tool on their own personal computers, which are not powerful enough to run the tool offline.

Constraint	Rationale
Usage of the Google Colab platform for training shall not exceed 100 compute units.	We will only purchase 100 compute units to stay within our project budget.
The training and tuning of our language models will be written in python.	Google Colab only supports python.
The number of training iterations of our language model will not exceed 1000.	The amount of computational power provided by Google Colab is the limiting factor in our training speed.

3.4 Off-the-Shelf Software

Constraint	Rationale
Any third-party code used in the project must be licensed under an open-source license that is compatible with the GPLv3 License.	Our project uses the GPLv3 license.

3.5 Anticipated Workplace Environment

There are no constraints imposed upon us by the anticipated workplace environment.

3.6 Schedule Constraints

Constraint	Rationale
Training data for language models shall consist of pre-labelled datasets on the internet, or scraped from online programming contests results and labelled with an existing plagiarism detection tool.	Given the schedule of the project, there is not enough time to source and label datasets by hand.

3.7 Budget Constraints

Constraint	Rationale
The cost of sourcing training data and training NLP models using cloud computers shall not exceed \$150 CAD.	We are students and we don't have a lot of money.

3.8 Enterprise Constraints

There are no enterprise constraints imposed upon our project.

4 Naming Conventions and Terminology

4.1 Glossary of All Terms, Including Acronyms, Used by Stakeholders involved in the Project

Insert your content here.

5 Relevant Facts And Assumptions

5.1 Relevant Facts

- The current standard for code plagiarism detection, MOSS, primarily relies on token matching and syntax-based comparison. This method lacks the ability to detect deeper semantic similarities in code.
- NLP techniques have advanced significantly in recent years, enabling more accurate natural language understanding. These techniques can be adapted to understand the structure and semantics of code, which could enhance plagiarism detection systems.
- There is a growing need for a plagiarism detection system that accounts for sophisticated plagiarism techniques, such as variable renaming, code restructuring, and adding non-functional code.
- Academic institutions are increasingly concerned with the fairness and accuracy of plagiarism detection systems to avoid penalizing students unfairly, especially with the rising prevalence of online and remote learning.

5.2 Business Rules

- The system must ensure compliance with data protection regulations by following a zero data retention policy.
- The similarity threshold for flagging plagiarism should be customizable by the institution or professor, allowing flexibility based on course policies.
- False positives (e.g., common code patterns) should be minimized, with options for professors to override flagged instances and manually validate the results.
- The system must be scalable to accommodate large datasets and multiple users submitting code for comparison at the same time.

5.3 Assumptions

- It is assumed that the academic institutions adopting this system have clear plagiarism policies and can provide a threshold score that reflects their definitions of plagiarism.
- It is assumed that the code samples provided for comparison are original and not previously processed by other plagiarism detection systems, ensuring that the results reflect real-time analysis.
- It is assumed that professors and administrators will review flagged cases manually to confirm plagiarism before taking any disciplinary action.
- It is assumed that students will not have access to the internal workings of the plagiarism detection algorithm, preventing them from finding potential loopholes to bypass detection.
- Software will be used only in Canada, and the legal and ethical considerations of this country will be taken into account during development.

6 The Scope of the Work

6.1 The Current Situation

The current code plagiarism detection tools such as MOSS rely on tokenization and syntax-level comparisons. Although effective in detecting direct copies or slight variations, they struggle when faced with techniques such as adding redundant code which allows the user to completely bypass detection while still plagiarizing the underlying logic and structure of the code.

Additionally, MOSS does not take into account for the complexity or intent behind the code, leading to issues such as false positives for common programming patterns. This creates a gap for more advanced tools capable of understanding the semantic meaning of code to more accurately detect plagiarism.

6.2 The Context of the Work

Our project aims to address these gaps by incorporating Natural Language Processing (NLP) and machine learning techniques which will be leveraged

to improve the accuracy of detecting copied code. The context of the work is within academic institutions, where the integrity of student work is paramount and our tool will be used by professors to ensure a fair grading process while also supporting students in understanding the ethical use of code.

6.3 Work Partitioning

- **Research and Design:** Conduct research on current plagiarism detection systems and state-of-the-art NLP techniques applicable to code plagiarism.
- **Data Collection:** Gather a dataset of code snippets, including both plagiarized and original works, to train and test the model.
- **Model Development:** Develop the NLP-based model capable of understanding the semantic meaning of code. This may involve exploring techniques like abstract syntax trees (ASTs), vector embeddings, or other representations of code that retain semantic meaning.
- **System Integration:** Build the system to take code as input, run through the developed model, and output a similarity score with appropriate thresholds.
- **Testing and Validation:** Test the system with various code samples to validate its performance and accuracy compared to traditional systems like MOSS. This will also test whether our method produces any false positives.
- **Documentation and Deployment:** Document the system architecture, the model, and the results. Deploy the system for use within academic settings.

6.4 Specifying a Business Use Case (BUC)

Business Use Case: Automated Code Plagiarism Detection for Academic Institutions

- **Actors:** Professors, Students, System Administrators

- **Trigger:** A professor or system administrator uploads multiple code submissions for plagiarism detection in a course assignment.

Main Success Scenario

1. The system ingests the uploaded code submissions.
2. The system processes each code snippet using the NLP model to generate semantic representations of the code.
3. The system compares the representations to detect plagiarism, taking into account code similarity beyond syntax or token matching.
4. The system outputs a similarity score for each comparison, with thresholds indicating whether plagiarism is suspected.
5. The professor reviews the similarity scores and flags any suspicious cases for further investigation.
6. The system generates a report summarizing the findings for the professor's review.

Extensions

- If the system detects false positives (common programming patterns being flagged as plagiarism), the professor can override the result.
- If new sophisticated plagiarism techniques are detected, the system can update its learning algorithms to improve accuracy over time.

7 Business Data Model and Data Dictionary

7.1 Business Data Model

Insert your content here.

7.2 Data Dictionary

Insert your content here.

8 The Scope of the Product

8.1 Product Boundary

Insert your content here.

8.2 Product Use Case Table

Insert your content here.

8.3 Individual Product Use Cases (PUC's)

Insert your content here.

9 Functional Requirements

9.1 Functional Requirements

- The system will take in code snippets as inputs
- The system must parse the input into a format that can be fed into a model for interpretation
- The system will leverage known natural language processing techniques to handle inputs
- The system will process the formatted input to provide a verdict on the presence of plagiarism as output
- The system must generate clear and concise outputs to indicate plagiarism
- The system's outputs should provide explanations on how to interpret output
- The system will allow users to create an account with an email and password
- The system will allow users to sign into their account with their email and password

- The system will send an email of the results in a .zip folder to a designated email address
- Users will be able to upload .zip folders of previous results to display the results on the interface

10 Look and Feel Requirements

10.1 Appearance Requirements

The user interface (UI) must adhere to the following appearance guidelines:

- **Consistency:** The UI should maintain a uniform color palette, font, and layout across all screens and elements. For instance, the primary color is defined as blue (#0047AB) for buttons, links, and headers, and white (#FFFFFF) for background areas.
- **Clarity:** All icons, buttons, and menus should be intuitive and clearly identifiable. Hovering over icons will provide tooltips with a brief description of their functionality.
- **Responsiveness:** The layout should adjust according to screen size, ensuring the interface remains usable on a range of devices, including desktops, tablets, and mobile phones.

10.2 Style Requirements

The style guidelines for the interface are as follows:

- **Typography:** The font family used across the UI will be *Roboto*. Headers should use a 24px font size, body text should be 14px, and button text should be 16px.
- **Color Scheme:** Buttons and interactive elements should use the primary color #0047AB. Background areas will use #F0F0F0, and error messages will be highlighted in #FF0000.
- **Button Styles:** All buttons should have rounded corners with a radius of 5px. On hover, the button background will lighten by 20%.

- **Spacing and Padding:** There should be at least 10px of padding between elements and a margin of 20px around each section to maintain a clean layout.

11 Usability and Humanity Requirements

11.1 Ease of Use Requirements

The system should be intuitive and simple to use for the target audience.

- **Minimal Learning Curve:** Users should be able to complete tasks with minimal instruction or training, taking at most 5 minutes from start to finish. The interface must guide users intuitively through this process.
- **Task Efficiency:** Common tasks should be achievable in no more than three clicks or interactions from the main screen.
- **Clear Navigation:** All navigation elements should be labeled clearly and positioned consistently across different pages to avoid confusion.

11.2 Personalization and Internationalization Requirements

The system will support a customizable experience in the future, but users will not be able to modify key settings such as themes or layout preferences at launch.

- **Future Customization Support:** Future updates will allow users to modify settings such as themes and layouts to suit their personal preferences.
- **English Language Support:** The system will operate exclusively in English, with all dates, currency, and numeric formats adhering to English (US) standards.

11.3 Learning Requirements

The system should provide clear documentation and onboarding materials.

- **Onboarding:** A guided onboarding process should be available for new users, helping them understand the main features of the system in under 5 minutes.
- **Help Documentation:** Detailed help documentation and tooltips should be available for key features to reduce the need for external assistance.

11.4 Understandability and Politeness Requirements

The system should use clear language and maintain a polite tone in all interactions with users.

- **Clear Language:** All messages and labels should be in simple, every-day language to ensure clarity for users of different levels of technical expertise.
- **Politeness:** Error messages and prompts should be worded politely and offer constructive guidance. Example error messages include:
 - *“Oops! Something went wrong. Please try again or contact support if the issue persists.”*
 - *“We’re sorry, but the file you uploaded is not supported. Please upload a .py file.”*

11.5 Accessibility Requirements

Due to time constraints, the system will not be fully accessible at launch. But this can be a future goal.

12 Performance Requirements

12.1 Speed and Latency Requirements

- The system should process the inputs and provide results in under one minute

- The system must notify users if processing surpasses one minute

12.2 Safety-Critical Requirements

- The system must avoid false positives to protect students from wrongful accusations.
- A manual override option for professors must exist to prevent automatic accusations solely based on the model.

12.3 Precision or Accuracy Requirements

- The system should have an accuracy rate of at least 90% for identifying plagiarized content.
- The false positive rate (incorrectly flagged plagiarism cases) must be kept below 5%.

12.4 Robustness or Fault-Tolerance Requirements

- The system will not crash in the case of a malformed input, and will instead issue an error message to the user.

12.5 Capacity Requirements

- The system should be able to handle batches of inputs without significant delays.
- The system should be able to handle batches of inputs without the total processing time exceeding 60 seconds.

12.6 Scalability or Extensibility Requirements

- The system should be designed in a way such that adding support for new file types or coding languages should not impede with current functionality

12.7 Longevity Requirements

- Backwards compatibility must be maintained when new features are introduced
- The system should be designed such that the natural language processing part can be changed in accordance to new research in the field.

13 Operational and Environmental Requirements

13.1 Expected Physical Environment

Insert your content here.

13.2 Wider Environment Requirements

Insert your content here.

13.3 Requirements for Interfacing with Adjacent Systems

Insert your content here.

13.4 Productization Requirements

Insert your content here.

13.5 Release Requirements

Insert your content here.

14 Maintainability and Support Requirements

14.1 Maintenance Requirements

- **Versioning:** keep track of and store model versions to allow switches between experimental models and previous versions (metric is version

history exists).

Justification: Gives ability to work on separate model architectures to try improvements in different directions as well as rollback on model releases if something goes wrong as well as provide experimental releases.

- **Metric Reports:** every model release shall be accompanied by a report of metrics decided upon to track model performance (metric is report exists).

Justification: Gives clarity whether model is being improved or not and what specific changes may be damaging the model. Will allow model health and performance to be tracked and maintained over time.

- **Issue Tracking:** every bug should be registered in an issue tracking system (metric is bugs existing as issues).

Justification: Brings attention to what bugs have been resolved and which ones still require work.

14.2 Supportability Requirements

- **Documentation:** provide concise (max 30 lines) comments on all apis and algorithms used in code base. This will allow others to comprehend what sections may need updates or bug fixes (metric is comment exists at every piece of api code or algorithm)

Justification: leaves no obfuscation in where somebody should adjust code or update a component by clarifying what responsibility each area of code holds.

- **Logging :** provides a report of code components that executed leading up to a crash whenever a crash occurs, so bugs can be identified and fixes can be staged. (metric is a log with relevant information about crash, such as a stack trace, is generated for user to view when crash occurs).

Justification: Almost 100% necessary for effectively debugging issues in the code and giving the ability to support the model in the future for people who aren't familiar with the code base.

- **Community Acknowledgement:** provide pathway for users to post or vote for requests/issues that should be addressed. (Metric is existence of said pathway, like email or github issues or forum).

Justification: Will highlight to community what are current pain points that should be prioritized and communication about how it is desired to address them.

14.3 Adaptability Requirements

- **Dataset Compatibility:** provide compatibility for at least two data formats when it comes to training datasets (metric is data formats that are accepted).

Justification: will allow model to accommodate wider amount of datasets and provides it better longevity in case either one of its data formats becomes less popular.

- **Template Adherence:** every model layer/component follows a template. (metric is a template exists for all model layers/components and are upheld by said layers/components.)

Justification - if somebody wishes to update the model with new layers/components or modify existing ones, there will be a guide to make the process straightforward and encourage those conducting the update to follow through.

- **Modularity:** every model layer/component can be treated as an individual function. (metric is model layers/components can be fed an input and produce an output which reflects only operations conducted within the layer/component).

Justification - enables removal or substitution of individual model layers/components which may be problematic or less efficient, similar to how one may remove or substitute a function in a program, without having to overhaul a larger/more significant portion of the model architecture.

15 Security Requirements

15.1 Access Requirements

- **Authentication:** The system shall require user authentication before allowing access to the tool. Only authenticated users can upload code for comparison.

Justification - Users need to be identifiable so that they can receive their results via email and so that usage limits can be imposed upon them.

15.2 Integrity Requirements

Integrity requirements are outside the scope of this project. The system is designed to avoid storing any user data. Since no data is retained, data integrity is not a concern.

15.3 Privacy Requirements

- **Zero Data Retention:** The system shall retain no user-uploaded data beyond the immediate need of the task. Refer to section 17.2.5 for more information.

Justification - This ensures that user privacy and protects sensitive academic work from being mishandled.

- **In Transit Encryption:** All data (including code snippets) shall be encrypted while being uploaded to or downloaded from the system using secure encryption protocols.

Justification - Protecting user data during transmission ensures that it can't be intercepted by unauthorized third parties.

15.4 Audit Requirements

- **In Transit Encryption:** The system shall log all user interactions and actions (such as authentication events, code uploads, and plagiarism checks) anonymously, without storing any user-uploaded code or sensitive data.

Justification - While no user data is retained, logging system events and user actions provides a transparent way to track system usage and ensure accountability without violating the zero data retention policy.

15.5 Immunity Requirements

Immunity requirements are beyond the scope of this project. The system should not need to defend against external attacks beyond the security measures already in place (e.g., authentication, encryption). The focus of the project is on code plagiarism detection, and immunity against malicious actors is not a primary concern.

16 Cultural Requirements

16.1 Cultural Requirements

No major cultural requirements are identified for this project but some that could be taken into consideration are:

Data Privacy and Ethical Use

Student Privacy: In some cultures and institutions, the handling of student work and data is highly regulated. Laws like the FIPPA mandate strict data privacy standards. The system should ensure that student data, including their code submissions, is securely handled, anonymized where possible, and not stored unnecessarily.

Differences in Academic Integrity Norms

Varying Definitions of Plagiarism: Some cultures and institutions promote collaboration as well as code borrowing so it is essential to define what plagiarism is in the context of this project. The tool should also be modifiable in its threshold for detecting plagiarism so institutions can change it to their needs.

17 Compliance Requirements

In developing the enhanced plagiarism detection tool, it is imperative to address various compliance requirements to ensure the tool operates legally,

ethically, and in alignment with industry standards. These requirements encompass legal obligations related to data protection, intellectual property rights, and adherence to educational policies, as well as compliance with established software development and data security standards.

17.1 Legal Requirements

1. **Data Protection and Privacy Laws:** The tool will process sensitive information, including students' code submissions, which may be considered personal data under Canadian privacy laws such as the *Personal Information Protection and Electronic Documents Act* (PIPEDA) at the federal level, and Ontario's *Freedom of Information and Protection of Privacy Act* (FIPPA) for public institutions. Compliance with these laws requires:
 - **Lawful Basis for Data Processing:** Ensuring that the collection and use of personal information is authorized under PIPEDA or FIPPA, typically requiring consent from students before processing their code or ensuring that processing is necessary for educational purposes.
 - **Data Minimization and Purpose Limitation:** Collecting only the data necessary for plagiarism detection and using it solely for that purpose.
 - **Transparency and Information Rights:** Informing students about how their data will be used, stored, and protected, and respecting their rights to access, correct, or withdraw their personal information.
 - **Security Measures:** Implementing appropriate technical and organizational measures to safeguard personal data against unauthorized access, loss, or disclosure, as required under PIPEDA and FIPPA.
2. **Intellectual Property Rights:** Under the *Copyright Act* of Canada, students typically hold the intellectual property rights to their original code. The tool must:
 - **Respect Ownership:** Use students' code exclusively for plagiarism detection without unauthorized distribution or reproduction.

- **Establish Clear Terms:** Provide clear terms of service or agreements outlining how the code will be used, ensuring students are aware and consent to these terms.
 - **Avoid Infringement:** Ensure that any storage or processing of code does not violate the *Copyright Act* or institutional policies.
3. **Academic Integrity Policies:** The tool must align with the academic integrity and misconduct policies of Canadian educational institutions by:
- **Supporting Fair Evaluation:** Assisting educators in identifying potential plagiarism accurately without bias.
 - **Due Process:** Ensuring that students have the opportunity to respond to plagiarism accusations, with results from the tool serving as part of a broader investigation rather than definitive proof.
 - **Confidentiality:** Maintaining the confidentiality of students' work and any findings related to plagiarism investigations.

17.2 Standards Compliance Requirements

1. **Software Development Standards:** Adherence to recognized software development practices is essential for ensuring quality and maintainability.
 - **SOLID Principles:** The project will follow SOLID principles—Single Responsibility, Open/Closed, Liskov Substitution, Interface Segregation, and Dependency Inversion—to promote clean, maintainable, and scalable code.
 - **Documentation and Testing:** Thorough documentation will be maintained throughout development, and rigorous testing will be conducted to validate the tool's performance and reliability.
2. **Data Security Standards:** Protecting sensitive data is a priority, and although full compliance with industry security standards such as ISO/IEC 27001 may not be feasible for a student capstone project, we will take measures to ensure data security.

- **OWASP Guidelines:** Basic security measures will be implemented in line with the Open Web Application Security Project (OWASP) guidelines to mitigate common vulnerabilities such as injection attacks and unauthorized access.
3. **Accessibility Standards:** The tool will be designed to support basic accessibility features.
 - **Screen Reader Support:** Key elements of the user interface will be made compatible with screen readers to assist visually impaired users.
 - **Alt Text for Images:** All images and non-text content will include descriptive alt text to improve accessibility for users relying on assistive technologies.
 - **Text Color Contrast:** The tool will ensure sufficient contrast between text and background colors to improve readability for users with visual impairments or color blindness.
 4. **Ethical AI and Machine Learning Standards:** As the tool leverages AI technologies, it must adhere to ethical standards in AI development.
 - **Transparency and Explainability:** Ensuring that the AI models used are transparent in their operation and that their decision-making processes can be explained to users.
 - **Fairness and Non-Discrimination:** Preventing biases in the AI models that could unfairly target or disadvantage any group of students. Ensuring that names and other potentially discriminatory information are not used in the detection process.
 5. **Data Handling and Retention Policies:** Establishing clear policies for how data is managed throughout its lifecycle.
 - **Zero Data Retention:** No user data will be retained beyond the immediate needs of the task. All user-uploaded data will be immediately loaded into memory and subsequently removed from disk as soon as the task is finished.
 - **Secure Disposal:** Implementing procedures for the secure deletion or anonymization of data that is no longer needed.

- **Audit and Compliance:** Regularly auditing data handling practices to ensure ongoing compliance with all relevant laws and standards.

By meticulously addressing these legal and standards compliance requirements, the project not only safeguards the rights and interests of all stakeholders but also enhances the credibility and trustworthiness of the plagiarism detection tool. Ensuring compliance is fundamental to the tool's success and its acceptance by educational institutions, educators, and students alike.

18 Open Issues

18.1 System Design Uncertainty

- How the integration of the Natural Language Processing (NLP) model with the code submission system will work. Such as choosing what NLP model to use, how to train it, and where to host it.
- Selecting the most appropriate techniques for semantic code analysis, such as Abstract Syntax Trees (ASTs) or vector embeddings.
- Deciding on how the model will handle false positives, especially common coding patterns, to avoid misclassification. Example. students using the same code snippets from a tutorial.

18.2 Data Collection Challenges

- While the project mentions scraping datasets from online programming contests and using pre-labeled datasets, there could be issues with data availability, quality, or ensuring that the data accurately reflects real-world usage.

18.3 Scalability Concerns

- Ensuring the system remains performant when processing large batches of submissions within the 10-minute processing time goal. Handling larger academic institutions with potentially thousands of submissions may challenge the current system's capacity requirements

- Given the limited budget for cloud compute units, it may be challenging to balance performance improvements with cost constraints. If training processes requires more resources than anticipated, costs could quickly escalate.

19 Off-the-Shelf Solutions

19.1 Ready-Made Products

There are many ready-made tools that aim to detect source-code level plagiarism, the most well-known of which being [MOSS](#). MOSS is most commonly used by professors. There are also a handful of open-source alternatives, including: [JPlag](#), [SIM](#), [Sherlock](#), and [Plaggie](#). Most of them support checking of multiple languages and use a variety of techniques to improve detection rates. These tools are relatively old and are not all actively developed.

19.2 Reusable Components

JPlag uses ANTLR 4 as a parser generator for many of its supported languages. For cross-language support, our tool can reuse JPlag's ANTLR grammar files to create language frontend parsers for each language we choose to support. Using a parser generator with pre-existing grammar files would reduce development time significantly, since the alternative would entail writing a custom parser for each supported language. By using ANTLR and JPlag's grammar files, we could feasibly support many source languages as opposed to just one, which would most likely be python (our product will be written in python, and python is capable of parsing it's own syntax tree).

19.3 Products That Can Be Copied

The primary inspiration for our product is MOSS. Our product, similar to many others, will copy the general data pipeline of input source code. Specifically, after reading an input source code file, plagiarism checkers typically have parsing, tokenizing, and normalizing steps. This is followed by some analysis on the normalized text - [MOSS uses "Winnowing"](#), an algorithm that produces local fingerprints in a piece of text.

To reduce development time, we plan to copy the first step (the text preprocessing as described above) of the data pipeline implemented in MOSS. This kind of text normalization is commonly studied and there are many resources that explain implementation details.

20 New Problems

20.1 Effects on the Current Environment

- **Blind Trust in the Tool:** One potential issue is that users may develop a sense of over-reliance on the tool based solely on the fact that it is AI-based. This can lead to blind trust in its results. This could cause professors to forego manual checks, resulting in incorrect classifications.

20.2 Effects on the Installed Systems

There should be no new problems for users of this tool since they don't need to install any software on their local machines. The tool's interface is accessed by a web browser, and the processing happens on the server side.

20.3 Potential User Problems

- **Difficulty with the New UI:** Some users may struggle to navigate the new UI, resulting in misuse of the tool. Additionally, some settings may not be known by the user and left at their default value, which may cause confusion or result in worse detection rate.
- **Misinterpretation of Results:** Users may incorrectly interpret the significance of the results provided by the tool. For instance, a high similarity score might be misinterpreted as definitive proof of plagiarism when there could be legitimate reasons for the similarity, such as the use of common libraries or design patterns.

20.4 Limitations in the Anticipated Implementation Environment That May Inhibit the New Product

- **Limited Compute Power Causing Delays:** If the system is deployed in environments with limited computational resources, such as local servers with insufficient processing power or memory, users may experience longer waiting times for results. Additionally, high computational demands may overload the system during peak usage times, also resulting in longer waiting times.

20.5 Follow-Up Problems

- **Server Maintenance and Costs:** Long-term maintenance of the server infrastructure could present challenges, including server downtime, performance degradation, and increasing operational costs as more users adopt the system. Costs related to server hosting, cloud processing, and ensuring uptime can grow as demand increases.
- **Model Upgrades and Performance Regressions:** Updating or training new models may introduce performance regressions or unexpected behavior in the plagiarism detection system. For example, a newer model may detect plagiarism more accurately but at the cost of higher processing times or an increase in false positives.

21 Tasks

21.1 Project Planning

1. Hazard Analysis (October 12 - October 23)

- **High-Level Description:** This document identifies potential risks or hazards in the project, assesses their severity, and proposes mitigations. It is important for ensuring safety and compliance with standards.
- **Story Points:** 12 days
- **Tasks:** Brainstorming hazards, analysis, and review.

2. V&V Plan (Revision 0) (October 24 - November 1)

- **High-Level Description:** The Verification and Validation (V&V) plan outlines the strategy for ensuring that the system is both built correctly and performs as intended. It covers the methods and resources required for testing.
- **Story Points:** 7 days
- **Tasks:** Define tests, acceptance criteria, and procedures for verification and validation.

3. Prep for Proof of Concept Demonstration (November 2 – November 11)

- **High-Level Description:** Preparation for a demonstration showing that the core concepts of the system are feasible. It typically involves implementing and testing essential features and preparing to showcase them to stakeholders.
- **Story Points:** 10 day
- **Tasks:** Development of key features and preparation for the demo.

4. Proof of Concept Demonstration (TBD)

- **High-Level Description:** A demonstration showing that the core concepts of the system are feasible. This involves implementing and testing essential features and demonstrating them to stakeholders.
- **Story Points:** 1 day
- **Tasks:** Execution and presentation of the demo.

5. Design Document (Revision 0) (November 23 - January 15)

- **High-Level Description:** This document outlines the architecture, components, data flow, and design decisions for the system. It serves as a blueprint for the development team.
- **Story Points:** 7 days
- **Tasks:** Detailed system design, architecture planning, and review.

6. **Prep for Revision 0 Demonstration (January 16 – February 3)**

- **High-Level Description:** Preparation for the Revision 0 Demo, which demonstrates functionality that meets the initial design and requirements.
- **Story Points:** 8 days
- **Tasks:** Prepare, refine, and finalize core product features.

7. **Revision 0 Demonstration (Date: TBD)**

- **High-Level Description:** A demonstration of the system after the first iteration, showing functionality that meets the initial design and requirements.
- **Story Points:** 1 day
- **Tasks:** Present the system, showcase its core features.

8. **V&V Report (Revision 0) (February 15 - March 7)**

- **High-Level Description:** This documents the process of verifying that the system meets its functional and nonfunctional requirements and validating that the system performs as intended. It includes the evaluation of both functional and nonfunctional requirements, details of unit testing, automated testing, and code coverage metrics. The report also highlights any changes made due to testing and traces these changes back to both requirements and modules. Comparisons to existing implementations are provided, where applicable.
- **Story Points:** 5 days
- **Tasks:** Collect results, analyze tests, and create the report.

9. **Final Demonstration (Revision 1) (Date: TBD)**

- **High-Level Description:** A final demonstration showcasing the complete system. By this stage, the system should be fully functional, meeting all requirements and passing validation.
- **Story Points:** 10 days

- **Tasks:** Ensure all components work seamlessly and prepare the final demo.

10. EXPO Demonstration (April TBD)

- **High-Level Description:** This is a public or stakeholder-facing presentation of the system, potentially at an event or exhibition. The focus is on clear communication and showcasing the system’s capabilities.
- **Story Points:** 7 days
- **Tasks:** Practice, setup, and running the demo smoothly.

11. Final Documentation (Revision 1) (Due: April 2)

- **High-Level Description:** This is the comprehensive final documentation, which includes all deliverables like the problem statement, development plan, POC plan, requirements document, design, and more.
- **Story Points:** 12 days
- **Tasks:** Complete, revise, and polish all components of the documentation.

21.2 Planning of the Development Phases

Phase 1: Proof of Concept (Oct 12 - Nov 11)

- **Model Development:**
 - Story Points: 30 days
 - Concentrate on training and evaluating the plagiarism detection model. Experiment with various algorithms and techniques.
 - Weekly Iterations: Regularly assess model performance, making adjustments based on results. Document findings and refine methods to enhance the model’s accuracy.

Phase 2: Initial System Design and UI Development (Nov 12 - Feb 3)

- **Model Development:**

- Story Points: 30 days
- Continue refining the model based on insights from the POC. Prepare the model for integration with the overall system architecture.
- Mid-January: Finalize the model for production use.

- **UI Development:**

- Begin designing the user interface. Focus on mockups and user experience workflows based on expected model outputs.
- Weekly Check-Ins: Collaborate with the model development team to ensure UI designs align with the model's capabilities and outputs.
- **Tasks and Duration:**

Task	Estimated Days	Notes
Account Creation UI	2	Interface for user registration
Comparison UI	20	Displaying similarity results
Report Generation UI	10	Generating and downloading reports

- **Unit and Integration Testing:**

- Story points: 5 days
- Conduct unit and integration tests between the finalized model and the backend components. Ensure data flow and output processing are functioning correctly.

Phase 3: Final Completion (Feb 4 - Mar 24)

- **Final Model Tweaks:**

- Make any last-minute adjustments to the model based on integration testing feedback.

- **Backend Development:**

- Develop the backend infrastructure, including user account management and result delivery systems.

- **Tasks and Duration:**

Task	Estimated Days	Notes
Account Management	1	Handling user registrations and login
Similarity Processing	14	Processing model outputs for similarity s
Report Generation	7	Generating reports based on model out
Report Delivery	4	Sending reports via email

- **UI Finalization:**

- Story points: 10 days
- Finalize the UI based on integration testing results and feedback from the model. Ensure all user functionalities are ready for testing.

- **Final Testing and Debugging:**

- Story points: 10 days
- Conduct comprehensive testing of the entire system, focusing on functionality, performance, and user experience. Address any bugs and finalize documentation.

22 Migration to the New Product

22.1 Requirements for Migration to the New Product

The migration to the new code plagiarism detection system must be carefully planned and executed to ensure smooth adoption by academic institutions. The following requirements should be addressed:

- **User Training and Support:** Provide comprehensive documentation and training materials for professors, system administrators, and students to familiarize them with the new system. This includes tutorials on how to upload code, interpret results, and resolve flagged cases.

- **Phased Rollout:** Implement a phased migration plan, starting with pilot tests in a controlled environment (e.g., one course or department) before full-scale implementation across the institution.
- **Data Security Compliance:** Ensure that all data migrations comply with data protection regulations, such as Canada’s Privacy Act.
- **System Downtime Minimization:** Plan the migration to minimize downtime and disruption to academic workflows. Ideally, the transition should occur during a break period, when student and faculty activity is low.

22.2 Data That Has to be Modified or Translated for the New System

For the migration to the new plagiarism detection system, certain data from the legacy systems must be modified or translated to ensure compatibility:

- **Code Submissions:** Legacy code submissions must be translated into a format that the new system can process, especially if the old system uses proprietary formats or different programming language encoding.
- **User Permissions:** Any user and administrator accounts need to be transferred to the new system. This includes roles, permissions, and access levels.
- **Configuration Data:** Settings from the old system, such as threshold scores, course configurations, and institution-specific policies, must be mapped and adjusted to fit the new system’s configuration parameters.
- **Metadata and Logs:** Metadata (e.g., anonymous submission timestamps, course IDs) and system logs related to prior plagiarism checks should be preserved and transferred, ensuring transparency and continuity.

23 Costs

The costs associated with this project come from several different parts of the project.

- Data must be acquired before training the model. The data required to train the model can cost money. However, the team intends to automate processes to acquire data, and thus there is no charge incurred. An unknown amount of money will be needed for data if the team's method of acquiring data fails.
- In the training and testing phase, the model will require hardware to be trained on. The team intends to use Google Co-lab and leverage the hardware provided by their cloud platform. This will cost approximately 30 dollars for the required computation. However, if more training and testing is required, more Google compute units will be used, costing more money. An upper limit of 150 dollars is set, which is 5 times more than the current guess.
- The front end will need to be hosted somewhere. However, free alternatives exist, thus this will have no cost incurred.

In total, the project should only cost approximately 30 dollars. The cost is subject to change, and can increase/decrease depending on the amount of data needed, and the how much Google's cloud hardware is used.

24 User Documentation and Training

24.1 User Documentation Requirements

Insert your content here.

24.2 Training Requirements

Insert your content here.

25 Waiting Room

Insert your content here.

26 Ideas for Solution

Insert your content here.

Appendix — Reflection

The information in this section will be used to evaluate the team members on the graduate attribute of Lifelong Learning. Please answer the following questions:

1. What knowledge and skills will the team collectively need to acquire to successfully complete this capstone project? Examples of possible knowledge to acquire include domain specific knowledge from the domain of your application, or software engineering knowledge, mechatronics knowledge or computer science knowledge. Skills may be related to technology, or writing, or presentation, or team management, etc. You should look to identify at least one item for each team member.
2. For each of the knowledge areas and skills identified in the previous question, what are at least two approaches to acquiring the knowledge or mastering the skill? Of the identified approaches, which will each team member pursue, and why did they make this choice?